

The neo-tavern memory system — a visual guide

A show-and-tell companion to the design record in [docs/memory.md](#). This file is all pictures + plain language: hand it to someone and they should be able to see how the whole thing works without reading a line of code. Every diagram is colour-coded to one legend (below).


Colour legend

Colour	Means
 Canon	The append-only message log — the source of truth, kept forever
 Segment	A <i>verbatim</i> 8-message block (raw transcript, embedded)
 Digest	A <i>distilled</i> structured summary of a block (anchor + facts + keywords, embedded, tiered)
 Model	A local/hosted model doing work (embedder, reranker, summarizer)
 Read path	Retrieval + what the user/model ultimately receives
 Prompt / infra	Prompt assembly, gating, plumbing

1. The one idea

Everything below exists to serve a single insight:


Canon is append-only and kept *forever*. The model's context window only sees the recent tail — older turns fall out of *visibility*, not existence. Memory is the regenerable, high-signal index that reaches back past the window and pulls the relevant past forward.



Memory is orthogonal to compaction. Compaction (native to the agent-SDK runner) compresses what's *inside* the window when it fills up. Memory is a *different axis* — it recovers what's already *fallen out*. Because of that, memory works in **all four provider modes**, including the OpenRouter chat-completions / responses modes that have no compaction at all.

2. The substrate — one boundary, two lenses


Every chat is sliced into fixed **8-message blocks**. Each completed block is captured through **two complementary lenses**, stored as two first-class, foreign-keyed, vector-indexed tables:



Why two? A raw 8-message chunk embeds *noisily* — every chunk looks vaguely similar, so similarity search returns mush. The digest distills the **load-bearing signal**, so it retrieves cleanly. We keep **both, permanently, linked** by the same `(chatId, blockIdx, seq-span)`: the digest is the sharp search key; the segment is the verbatim ground truth it points back to.

3. What a *digest* actually is

A tier-0 digest is not a paragraph summary. It's a **structured, retrieval-optimised** unit with three parts, produced under a strict prompt:




A real digest the pipeline produced (block 2 of a 222-message chat; one character renamed for sharing):

```
[Wyatt & Bess – Midnight Crossing, the Unleashed RX-78-2]
- Wyatt revealed he bought a Bandai PG Perfect Grade Unleashed RX-78-2 Gundam ($350, 3,500 pieces) as his first kit – deliberately, to make Bess come help him build it.
- Bess committed to building the Unleashed with him and declared it nonnegotiable.
- Wyatt started watching Mobile Suit Gundam SEED (the "wrong" entry point) on purpose to bait Bess into coming over.
- Bess crossed the street at 11:42 PM, walked in without knocking, and took the tablet to stop him – they settled on 08th MS Team instead.
keywords: Unleashed RX-78-2, Perfect Grade, $350 kit, 3,500 pieces, Gundam SEED,
          08th MS Team, midnight crossing, build-it-together, foreclosure house, Samui
```

Keywords like *"Unleashed RX-78-2"* or *"08th MS Team"* are exactly the kind of distinctive tokens that retrieve precisely — which a raw transcript chunk would bury.

4. The write path — how memory gets built

Both lenses are generated **after a turn completes** (fire-and-forget, never blocking the reply) *and* in bulk by the backfill script for imported history. Same functions, same result.




The summarizer is free-first: a local **Qwen3-4B-Instruct GGUF** (in-process, node-llama-cpp) when configured, otherwise a hosted **Claude Haiku** fallback over the existing chat-completions runner. The embedder/reranker run in-process on the homelab GPUs (onnxruntime CUDA).

It self-heals. A block is only re-digested if it's **missing or stale** — its seq-span changed, or a message in it was edited *after* the digest was written. Swipes/edits at the live tip never touch settled digests (that's what the 8-message verbatim window protects), and a fork lazily rebuilds only what diverged.

5. The tiering — keeping “the story so far” bounded

If we injected *every* tier-0 digest, a 1000-message chat would inject 60+ of them. Instead, blocks **consolidate upward**: every `fanOut = 4` lower-tier digests merge into one coarser digest, with a delta prompt (“*here are prior consolidations — do NOT repeat them*”). The arc stays compact as the chat grows.



The tier-1 anchor for blocks 0–3 of that same chat read: “[*Wyatt & Bess — Twelve Years Later: a Reunion, a Gundam Build, and Old Family History*]” — a genuine cross-block synthesis, not a concatenation.


This is what the **tiered** read mode (next section) consumes: a **bridge** of coarse high-tier digests for the distant past + fine tier-0 digests for the recent past — so the injected “story so far” stays at a roughly constant token budget no matter how long the chat runs.

6. Two scopes, two read paths

The **same substrate** serves two very different questions. This is the heart of the system.

The within-chat path is **flat query-driven RAG**, modeled on SillyTavern's native Vector Storage: the **last 2 messages** form the query, the most relevant tier-0 digests are retrieved (+ reranked), and injected.

mixC is the default; **mixA** (inject every tier-0) and **tiered** (the consolidation bridge) are opt-in "give me the whole arc" modes. The injected set is tiny (~3 digests ≈ 1k tokens), so it's window-independent — the 32k/200k context size only bounds how much *raw* history the runner sends, not memory.



	Within-chat injection	Cross-chat corpus search
Question	“What happened earlier in <i>this</i> chat?”	“Where, across <i>all</i> my chats, did X happen?”
Scope	one chat	the whole owner corpus
Match	exact in-process cosine (small, this chat)	global ANN (<code>vector_top_k</code>) + CSLS
Sources	tier-0 digests	all digests <i>and</i> all segments (hybrid)
Reranks	yes — mixC is the default	yes — one joint list
Output	text that fills <code>{{memory}}</code>	ranked hits → <code>seq</code> spans back to canon

CSLS corrects “hub” vectors — a few generic digests that sit suspiciously close to *everything* — by penalising them with a per-group hubness score, so true relevance wins.


Retrieval in action — an illustrative run (the discrimination mechanism is what matters, and is unchanged by block size): the within-chat path embeds the query and cosines over this chat’s tier-0 digests:

Query	Top digest surfaced	sim
“which Gundam kit did they decide to build together?”	block 2 — [Wyatt & Bess — Midnight Crossing] (names the Unleashed RX-78-2)	0.53
“the night she ran across the street to stop him watching the wrong show”	block 2 — [Wyatt & Bess — Midnight Crossing]	0.56
“the four-hour garage talk about building mobile suits”	block 1 — [Wyatt & Bess — Garage reunion]	0.62


Note the discrimination: the *general* “mobile suits” talk pulls **block 1**, while the *specific* build-night decision pulls **block 2** — two adjacent, topically-similar blocks stay cleanly separable because the anchors + keywords give each a distinct fingerprint. Raw-chunk embeddings, which all collapse toward the same “two people chatting” vector, can’t make that split.

7. Where the injected memory lands

Within-chat retrieval fills the `{{memory}}` marker, which lives in the **dynamic (cache-safe)** half of the system prompt — *after* the cache boundary — so the per-turn memory set never busts the cached static prefix.



8. The whole thing, end to end



9. Built today vs. enabled later

The substrate above is the hard part, and it's **built and validated**. A few features sit *on top* of it with zero rework required — the design deliberately doesn't preclude them:

✓ Built	→ Enabled later (substrate-ready)
Structured tier-0 digests + hierarchical consolidation	Trackers — a single entry that updates in place (relationship status, inventory, plot threads)
Verbatim <code>chat_segments</code> + the hybrid corpus search	Clips — user-pinned one-off facts
CSLS hubness on both tables	User-curated long-term promotion — hand-pick what persists
Edit/fork-aware lazy regeneration	Per-chat summarizer profiles

Because every digest is a **pure function of canon**, none of these risk corrupting the source of truth — they're just additional lenses over the same append-only log.

Knobs (all per-preset, surfaced in the UI later)

`blockSize` (8) · `verbatimWindow` (8) · `queryWindow` (2) · `mode` (default **mixC**; off / mixA / mixB / mixC / tiered) · `fanOut` (4) · `maxTier` (3) · `retrieveK` (4) · `rerankTo` (3) · `minScore` (0.25) · `keywordMatch` · `summarizer` (local-first Qwen-4B → hosted Haiku fallback).

Defaults re-derived from the real corpus for the 32k window ([docs/memory.md](#) §9). All three vector tables — `character_embeddings`, `chat_digests`, `chat_segments` — are owner-keyed + FK'd. Schema: `src/db/schema.ts`.